



Understanding Student Debt

Statistics and Machine Learning

Adrian Halarewicz
Springboard Capstone Project



Problem

- Many students blindly take on student debt.
- Students struggle to define a repayment plan.
- Will the student be able to repay the debt?



Solution

- Explore U.S. Census Income data to better understand future financial standing.
- Loan data from Lending Club can help predict if a borrower will default on a Loan.



Example - Case Study

THE CLIENT :

Age → **24**

Gender → **Male**

City → **Los Angeles, CA**


Degree/Industry → **Scientific and Technical Services**

Amount Borrowed → **\$40,000**

Min. Monthly Payment → **\$450**

The Student Debt Landscape

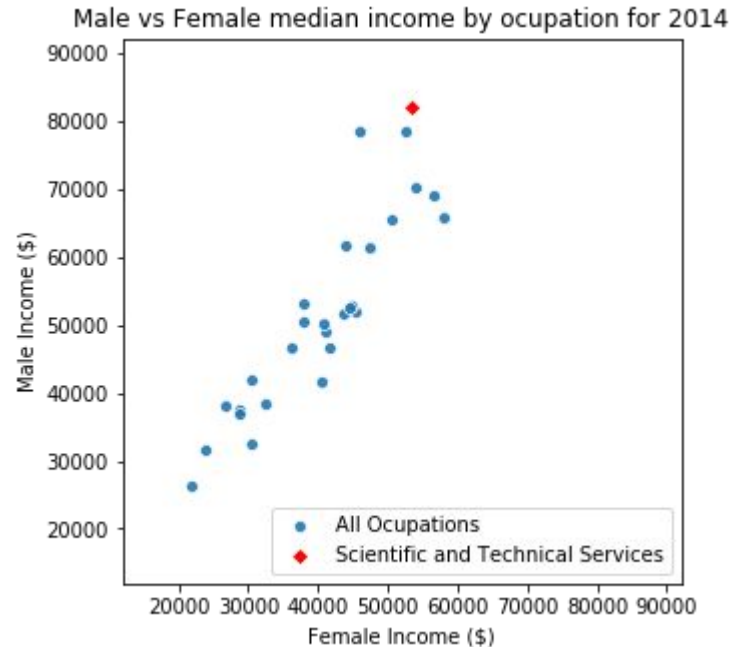
- Client Loan Balance: \$40,000
- More debt than most borrowers.



Median income by occupation


- Expected Income: \$83,000
- The median income for Science and Technical Services.
- Higher than most other occupational categories.
- A healthy Debt to Income Ratio:

Income : Debt → 83,000 : 40,000



How does the client compare?


- Median Income for 24 year olds:
 - \$34,000
- Mean Income for 24 year olds:
 - \$28,000



Client will be making significantly more
than other people his age.

Consider Cost of Living

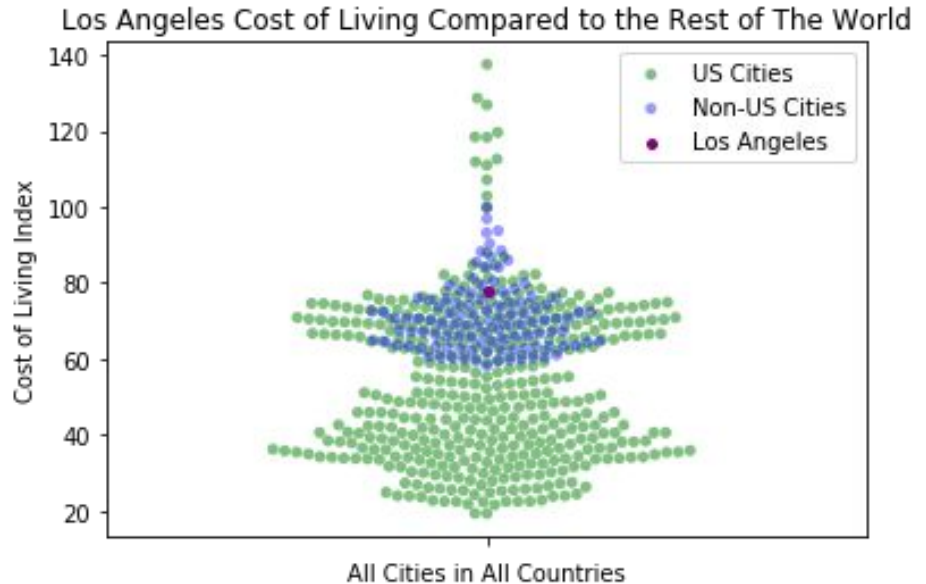
- Cost of Living in California is High!



- Should the client consider moving to a state with a Lower cost of Living?

Consider Cost of Living

- The Cost of Living in Los Angeles is among the most expensive cities in the U.S. and throughout the world.
- Fortunately, wages and salaries are typically higher in areas with a higher cost of living.
- Living in a high cost of living area is beneficial to the client's overall debt to income ratio.





Will the user Default on their Loan?

Implementing Machine Learning to Predict Loan Default

- Random Forest
 - An out of the box classifier for quickly making initial predictions with default hyperparameters.
- Support Vector Machine
 - Requires cross-validation for hyperparameter tuning.



Training Data

- Lending Club loan data: All loans from 2007 → 2019 Q3
- Tidy Data
 - Rows represent instances of loans issued by Lending Club.
 - Columns represent individual features describing each respective loan.
- 80,000+ rows of loans issued
 - a variety of purposes: Education, home, auto, debt consolidation, ...
 - Loan Status is listed for each loan.




Features

- 74 unique columns describing individual rows.
- Create dummy variable columns for categorical features.
 - `pd.get_dummies(column_name)`
- Convert discrete features to integers.
- Lasso Regression (Elastic Net):
 - Identify important features with most predictive power.
- Drop irrelevant columns.
- Final DataFrame contains 50 columns of features + 1 column for Loan Status = 51 total columns

Lasso Regression

Predictive Power:

- - loan_amnt
- + out_prncp
- - total_pymnt
- - total_rec_int
- + last_pymnt_amnt





Targets

- The Loan Status column:
 - 1 → Loan is in DEFAULT
 - 0 → Loan is NOT IN DEFAULT
- Fixing class imbalance
 - Minority class: Loans in Default
 - Majority class: Loans not in Default
 - Lending Club's best effort to prevent issuing loans that will default.



SMOTE for Class Imbalance

- Synthetic Minority Over-Sampling Technique
- Over-Sample the minority class.
- Create an equal number of instances
for Loans in Default.


```
# inspect class imbalance for defaulted loans
unique, count = np.unique(y_train, return_counts=True)
value_counts = {k:v for (k,v) in zip(unique, count)}
value_counts
```

```
{0: 619995, 1: 853}
```

```
# Apply Synthetic Minority Over-sampling Technique (SMOTE)
sm = SMOTE(random_state=42)
X_train_bal, y_train_bal = sm.fit_sample(X_train, y_train)
```

```
# inspect balanced training data
unique, count = np.unique(y_train_bal, return_counts=True)
value_counts = {k:v for (k,v) in zip(unique, count)}
value_counts
```

```
{0: 619995, 1: 619995}
```



The Training Data and Hold-Out Set

- Drop rows with missing data.
- Extract targets and features.
- Split the targets and features into train sets and test/hold-out sets.

```
loans = loans.dropna()  
loans.shape
```

```
(886927, 52)
```

```
y = loans['default'].values  
X = loans.drop('default', axis=1).values
```

```
X.shape
```

```
(886927, 51)
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=33)
```




Random Forest: Training the model

- Instantiate a Random Forest Classifier with Default hyperparameters.
- Fit the classifier to the training data.
- Evaluate the model using the hold-out set.

```
# instantiate random forest classifier
random_forest = RandomForestClassifier(max_depth=True)

# fit model to training data
print('fitting model to training data')
random_forest = random_forest.fit(X_train, y_train)

# test model performance
print('evaluating model performance')
score = random_forest.score(X_test, y_test)
print(score)
```

Random Forest: Evaluating the Model

- With Class Imbalance

- Accuracy = 0.9986 but TP = 0 FP = 0
- The model predicted that ALL LOANS WILL NOT DEFAULT
 - High accuracy due to severe class imbalance.

- After SMOTE

- Accuracy = 0.7132 but TP = 251 FP = 76180
- The model predicted that ALL LOANS WILL NOT DEFAULT.
 - Sacrifice Accuracy to better predict loan default.
 - True Positive Rate = 0.685

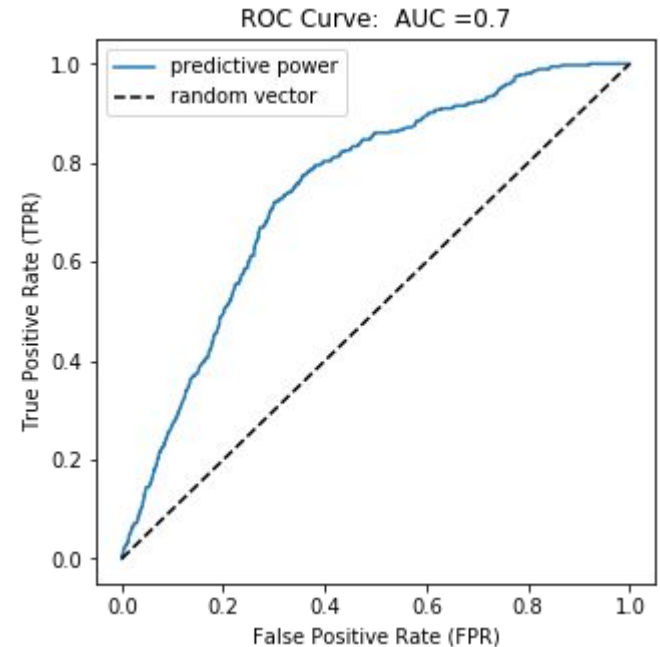
	Predict: NO	Predict: YES
Actual: NO	265713	0
Actual: YES	366	0

	Predict: NO	Predict: YES
Actual: NO	TN	FP
Actual: YES	FN	TP

	Predict: NO	Predict: YES
Actual: NO	189533	76180
Actual: YES	115	251

Random Forest: ROC Curve

- Area under the ROC Curve
 - AUC = 0.70
- This model has predictive power.
- Ability to make significant predictions.
- Outperforms a Random Vector.



Support Vector Machine: Training the model

- Instantiate the SVM classifier

- with default parameters

- Fit the model to the training data.

- Predict on the hold-out set.

- Before SMOTE:

- $TP = 0$ & $FP = 0$

- Apply smote to correct class imbalance in the training set and the model's predictions.

- NOTE: Unlike Random Forest, SVM requires the data to be scaled to prevent distance bias from occurring.

```
# instantiate SVM classifier  
svm = SVC()
```

```
# fit SVM to training data  
svm.fit(X_train, y_train)
```

```
SVC(C=1.0, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,  
    decision_function_shape='ovr', degree=3, gamma='scale', kernel='rbf',  
    max_iter=-1, probability=False, random_state=None, shrinking=True,  
    tol=0.001, verbose=False)
```

```
# make predictions on test features  
y_pred = svm.predict(X_test)
```

Support Vector Machine: GridSearchCV

- Hyperparameter tuning:
 - Grid Search with 5 fold Cross Validation.
 - Find best values of C and Gamma.
- $C \rightarrow$ the amount of slack given to outliers.
- γ \rightarrow the perpendicular distance of a point to the fit line in the opposite direction of \mathbf{w} .

```
# define hyperparameter space
c_values = [0.001, 0.01, 0.1, 1]
gamma_values = [0.001, 0.01, 0.1]
param_grid = {'C': c_values, 'gamma': gamma_values}
```

```
# create grid-search object
grid_search = GridSearchCV(clf, param_grid, cv=5)
```

```
# fit apply gridsearch to SVM with training data
grid_search.fit(X_train_bal, y_train_bal)
```

```
GridSearchCV(cv=5, error_score=nan,
              estimator=SVC(C=1.0, break_ties=False, cache_size=200,
                             class_weight=None, coef0=0.0,
                             decision_function_shape='ovr', degree=3,
                             gamma='scale', kernel='rbf', max_iter=-1,
                             probability=False, random_state=None, shrinking=True,
                             tol=0.001, verbose=False),
              iid='deprecated', n_jobs=None,
              param_grid={'C': [0.001, 0.01, 0.1, 1],
                           'gamma': [0.001, 0.01, 0.1]},
              pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
              scoring=None, verbose=0)
```

```
# identify best C and gamma combination
best_params = grid_search.best_params_
best_params
```

```
{'C': 1, 'gamma': 0.1}
```



Support Vector Machine: Making Predictions

- GridSearchCV stores and remembers the best combination of **C** and **gamma**.
- After the cross-validation predictions are made with the model trained with the best Hyperparameters.
- Large Data → Long training time for each model
 - Means cross-validation is an extremely long process.
 - This model was trained and evaluated on a sample of the data.
 - A Pipeline was built to:
 - Perform grid search with 5-fold cross validation.
 - Train the model on the complete data set with the best hyperparameters.
 - Make better predictions based on the complete data set.

```
# Run gridsearch cross validation  
cv = GridSearchCV(pipeline, param_grid=param_grid)  
cv.fit(X_train_bal, y_train_bal)  
  
# make predictions with best params from grid search  
y_pred = cv.predict(X_test)
```



Conclusion

Given the client's income and credit history, we can:

1. Illustrate where the student borrower's financial standing.
2. Predict if the client will default on their loan with 68% accuracy.

Random Forest is a great out of the box classifier.